

LANGUAGE-BASED INTERACTIVE TESTING

Yishuai Li

A DISSERTATION

in

Computer and Information Science

Presented to the Faculties of the University of Pennsylvania

in

Partial Fulfillment of the Requirements for the

Degree of Doctor of Philosophy

2022

Supervisor of Dissertation

Benjamin C. Pierce

Professor of Computer and Information Science

Graduate Group Chairperson

Mayur Naik, Professor of Computer and Information Science

Dissertation Committee

Steve Zdancewic, Professor of Computer and Information Science, Chair

Mayur Naik, Professor of Computer and Information Science

Boon Thau Loo, Professor of Computer and Information Science

John Hughes, Professor of Computing Science, Chalmers University of Technology

LANGUAGE-BASED INTERACTIVE TESTING

COPYRIGHT

2022

Yishuai Li

This work is licensed under a Creative Commons Attribution-ShareAlike 4.0 International (CC BY-SA 4.0) License

To view a copy of this license, visit

<https://creativecommons.org/licenses/by-sa/4.0/>

Acknowledgments

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

ABSTRACT

LANGUAGE-BASED INTERACTIVE TESTING

Yishuai Li

Benjamin C. Pierce

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

Contents

Title	i
Copyright	ii
Acknowledgments	iii
Abstract	iv
Contents	v
List of Figures	vi
Chapter 1. Introduction	1
1.1. Interactive Testing	1
1.2. Internal and external nondeterminism	2
1.3. Contribution	6
Introduce motivation in this section entirely	6
Chapter 2. Validator Theory	8
Chapter 3. Validator in Practice	11
3.1. Specification Languages	11
3.2. From Specification to Tester	12
Chapter 4. Test Harness Design	16
Chapter 5. Related Work	17
5.1. Specifying and Testing Protocols	17
5.2. Reasoning about Network Delays	18
Chapter 6. Discussions	19
Chapter 7. Conclusion	20
Bibliography	21
Appendix A. Unstructured contents	23
A.1. Challenges: Testing Internal and Network Nondeterminism	23
A.2. Specification Language	25
A.3. Derivation: from Server Specification to Testing Program	27
A.4. Evaluation	36

List of Figures

1.1	Ad hoc tester for HTTP/1.1 conditional requests.	5
3.1	Interaction trees and their traces of events.	11
3.2	Linear specification of the swap server. In the <code>linear_spec</code> ' loop, the parameter <code>conns</code> maintains the list of open connections, while <code>last_msg</code> holds the message received from the last client (which will be sent back to the next client). The server repeatedly chooses between accepting a new connection or doing a receive and then a send on some existing connection picked in the list <code>conns</code> . The linear specification is initialized with an empty set of connections and a message filled with zeros.	13
3.3	Symbolic model handling conditional PUT request. The model maintains two states: <code>data</code> that maps keys to their values, and <code>xtag</code> that maps keys to symbolic variables that represent their corresponding ETags. Upon receiving a PUT request conditioned over “If-Match: <code>t</code> ”, the server should decide whether the request ETag <code>matches</code> that stored in the server. Upon matching, the server processes the PUT request, and represents the updated value’s ETag as a fresh variable.	14
3.4	Network model for concurrent TCP connections. The model maintains a <code>buffer</code> of all packets en route. In each cycle, the model may nondeterministically branch to either absorb or emit a packet. Any absorbed packet is appended to the end of buffer. When emitting a packet, the model may choose a connection and send the oldest packet in it.	14
3.5	Deriving tester program from specification	15
A.1	Proxy model specifying a server that forwards a message immediately upon receiving it.	25
A.2	A reordered observation, with two valid network-level explanations.	25
A.3	Example of invalid observation.	26
A.4	Network model for concurrent TCP connections. The model maintains a <code>buffer</code> of all packets en route. In each cycle, the model may nondeterministically branch to either absorb or emit a packet. Any absorbed packet is appended to the end of buffer. When emitting a packet, the model may choose a connection and send the oldest packet in it.	27
A.5	Symbolic model handling conditional PUT request. The model maintains two states: <code>data</code> that maps keys to their values, and <code>xtag</code> that maps keys to symbolic variables that represent their corresponding ETags. Upon receiving	

- a PUT request conditioned over “If-Match: τ ”, the server should decide whether the request ETag **matches** that stored in the server. Upon matching, the server processes the PUT request, and represents the updated value’s ETag as a fresh variable. 28
- A.6 Deriving tester program from specification 28
- A.7 Interpretation example. **acc** receives a number and returns the sum of numbers received so far. **tee** prints all the numbers sent and received. Interpreting **acc** with interpreter **tee** results in a program that’s equivalent to **tee_acc**. 29
- A.8 Dualizing server model into observer model. Upon **recv** events, the observer generates a packet and sends it to the server. For **send** events, the observer receives a packet **p1**, and fails if it does not match the specified **pkt**. When the server makes nondeterministic **IF** branches, the observer **determines** between the branches by **unifying** the branch condition with its conjectured value, and then observing the corresponding branch. 30
- A.9 Instantiating symbolic events. The tester maintains a **unifyState** which stores the constraints on symbolic variables. When the specification creates a **fresh** symbol, the tester creates an entry for the symbol with no initial constraints. Upon **unify** and **guard** events, the tester checks whether the **assertion** is compatible with the current constraints. If yes, it updates the constraints and move on; otherwise, it raises an error on the current branch. 31
- A.10 From nondeterministic model to deterministic tester program. If the model makes nondeterministic branches, the tester picks a branch to start with, and puts the other branch into a set of other possibilities. If the current branch has failed, the tester looks for other possible branches to continue checking. When the current branch sends a packet, the tester filters the set of other possibilities, and only keeps the branches that match the current send event. If the model wants to receive a packet, the tester handles both cases whether some packet has arrived or not. 33
- A.11 Embedding programs’ internal state into the events. By expanding the events’ parameters, we enrich the test case generator’s knowledge along the interpretations. 34
- A.12 Composing **http** server model with **tcp** network model by interpreting their events and passing messages from one model to another. The composing function takes four parameters: server and network models as **srv** and **net**, and the message buffers between them. When **srv** wants to **send** a packet in Line 21, the packet is appended to the outgoing buffer **bo** until absorbed by **net** in Line 12, and eventually emitted to the client in Line 7. Conversely, packets sent by clients are absorbed by **net** in Line 13, emitted to the application’s incoming buffer **bi** in Line 6, until **srv** consumes it in Line 24. 35
- A.13 Cost of detecting bug in each server/mutant. The left box with median line is the tester’s execution time before rejecting the server, which includes interacting with the server and checking its responses. The right bar with

median circle is the number of HTTP/1.1 messages sent and received by the tester before finding the bug. Results beyond 25%–75% are covered by whiskers.

37

A.14 The trace on the left does not convince the tester that the server is buggy, because there exists a certain network delay that explains why the PUT request was not reflected in the 200 response. When the trace is ordered as shown on the right, the tester cannot imagine any network reordering that causes such observation, thus must reject the server.

38

CHAPTER 1

Introduction

We trust programs by testing them, but how do we trust our testers?

Software engineering requires rigorous testing of rapidly evolving programs, which costs manpower comparable to developing the product itself. To guarantee programs' compliance with the specification, we need testers that can tell compliant implementations from violating ones.

This thesis studies the testing of interactive systems' semantics: The system under test (SUT) interacts with the tester by sending and receiving messages, and the tester determines whether the messages sent by the SUT are valid or not with respect to the protocol specification.

This section introduces the basic concepts of interactive testing (Section 1.1), why nondeterminism made this problem difficult (Sections 1.2), and how language designs address the challenges introduced by nondeterminism (Section 1.3).

1.1. Interactive Testing

Suppose we want to test a web server that supports GET and PUT methods:

```
CoFixpoint server (data: key → value) :=  
  request ← recv;;  
  match request with  
  | GET k   ⇒ send (data k);; server data  
  | PUT k v ⇒ send Done   ;; server (data [k ↦ v])  
end.
```

We can write a tester client that interacts with the server and determines whether it behaves correctly:

```
CoFixpoint tester (data: key → value) :=  
  request ← random;;  
  send request;;  
  response ← recv;;  
  match request with  
  | GET k   ⇒ if response =? data k  
               then tester data  
               else reject  
  | PUT k v ⇒ if response =? Done  
               then tester (data [k ↦ v])  
               else reject  
end.
```

This tester implements a reference server internally that computes the expected behavior. The behavior is then compared against that produced by the SUT. The tester rejects the SUT upon any difference from the computed expectation.

The above tester can be viewed as two modules: (i) a *test harness* that interacts with the server and produces transactions of sends and receives, and (ii) a *validator* that determines whether the transactions are valid or not:

```

(* Compute the expected response and next state of the server. *)
Definition serverSpec request data :=
  match request with
  | GET k    => (data k, data)
  | PUT k v => (Done  , data [k ↦ v])
  end.

(* Validate the transaction against the stateful specification. *)
Definition validate spec request response data :=
  let (expect, next) := spec request data in
  if response ==? expect then Success next else Failure.

(* Produce transactions for the validator. *)
CoFixpoint harness validator state :=
  request ← random;;
  send request;;
  response ← recv;;
  if validator request response state is Success next
  then harness validator next
  else reject.

Definition tester := harness (validate serverSpec).

```

Such testing method works for deterministic systems, whose behavior can be precisely computed from its input. Whereas, many systems are allowed to behave nondeterministically. How to test systems that involve randomness? How to validate servers' behavior against concurrent clients? The following sections discuss nondeterminism by partitioning it in two ways, and explains how they pose challenges to the validator and the test harness.

1.2. Internal and external nondeterminism

When people talk to each other, voice is transmitted over substances. When testers interact with the SUT, messages are transmitted via the runtime environment. The specification might allow SUTs to behave differently from each other, just like people speaking in different accents, we call it *internal nondeterminism*. The runtime environment might affect the transmission of messages, just like solids transmit voice faster than liquids and gases, we call it *external nondeterminism*.

1.2.1. Internal nondeterminism. Within the SUT, correct behavior may be underspecified. For example, HTTP [7] allows requests to be conditional: If the client has a local copy of some resource and the copy on the server has not changed, then the server needn't resend the resource. To achieve this, an HTTP server may generate

a short string, called an “entity tag” (ETag), identifying the content of some resource, and send it to the client:

<i>/* Client: */</i>	<i>/* Server: */</i>
GET /target HTTP/1.1	HTTP/1.1 200 OK
	ETag: "tag-foo"
	... content of /target ...

The next time the client requests the same resource, it can include the ETag in the GET request, informing the server not to send the content if its ETag still matches:

<i>/* Client: */</i>	<i>/* Server: */</i>
GET /target HTTP/1.1	HTTP/1.1 304 Not Modified
If-None-Match: "tag-foo"	

If the ETag does not match, the server responds with code 200 and the updated content as usual.

Similarly, if a client wants to modify the server’s resource atomically by compare-and-swap, it can include the ETag in the PUT request as *If-Match* precondition, which instructs the server to only update the content if its current ETag matches:

<i>/* Client: */</i>	<i>/* Server: */</i>
PUT /target HTTP/1.1	HTTP/1.1 204 No Content
If-Match: "tag-foo"	
... content (A) ...	
 <i>/* Client: */</i>	 <i>/* Server: */</i>
GET /target HTTP/1.1	HTTP/1.1 200 OK
	ETag: "tag-bar"
	... content (A) ...

If the ETag does not match, then the server should not perform the requested operation, and should reject with code 412:

<i>/* Client: */</i>	<i>/* Server: */</i>
PUT /target HTTP/1.1	HTTP/1.1 412 Precondition Failed
If-Match: "tag-baz"	
... content (B) ...	

<pre> /* Client: */ GET /target HTTP/1.1 </pre>	<pre> /* Server: */ HTTP/1.1 200 ok ETag: "tag-bar" ... content (A) ... </pre>
---	--

Whether a server’s response should be judged *valid* or not depends on the ETag it generated when creating the resource. If the tester doesn’t know the server’s internal state (*e.g.*, before receiving any 200 response that includes an ETag), and cannot enumerate all of them (as ETags can be arbitrary strings), then it needs to maintain a space of all possible values, and narrow the space upon further interactions with the server. For example, “If the server has revealed some resource’s ETag as `"tag-foo"`, then it must not reject requests targetting this resource conditioned over `If-Match: "tag-foo"`, until the resource has been modified”; and “Had the server previously rejected an `If-Match` request, it must reject the same request until its target has been modified.”

This idea of remembering matched and mismatched ETags is implemented in Figure 1.1. For each key, the validator maintains three internal states: (i) The value stored in `data`, (ii) the corresponding resource’s ETag, if known by the tester, stored in `tag_is`, and (iii) ETags that should not match with the resource’s, stored in `tag_is_not`. Each pair of request and response contributes to the validator’s knowledge of the target resource. The tester rejects the SUT if the observed behavior does not match its knowledge gained in previous interactions.

Even a simple nondeterminism like ETags requires such careful design of the validator, based on thorough comprehension of the specification. For more complex protocols, we hope to construct the validator in a reasonable way.

1.2.2. External nondeterminism. Beyond the server, messages and responses between the server and different clients might be delayed and reordered by the network and operating-system buffering. If the tester cannot control how the execution environment reorders messages—*e.g.*, when testing over the Internet—it needs to specify what servers are valid as observed over the network.

These sources of nondeterminism pose challenges in various aspects of testing network protocols: (i) The *validation logic* should accept various implementations, as long as the behavior is included in the specification’s space of uncertainties; (ii) To capture bugs effectively, the *test harness* should generate test cases based on runtime observations; (iii) When *shrinking* a counterexample, the test harness should adjust the test cases based on the server’s behavior, which might vary from one execution to another.

To address these challenges, I introduce symbolic languages for writing specifications and representing test cases:

```

Definition validate request response
  (data      : key → value)
  (tag_is    : key → Maybe etag)
  (tag_is_not: key → list etag) :=
match request, response with
| PUT k t v, NoContent ⇒
  if t ∈ tag_is_not k then Failure
  else if (tag_is k =? Unknown) || strong_match (tag_is k) t
  then (* Now the tester knows that the data in [k]
        * is updated to [v], but its new ETag is unknown. *)
        Success (data      [k ↦ v],
                  tag_is    [k ↦ Unknown],
                  tag_is_not [k ↦ [] ])
  else Failure
| PUT k t v, PreconditionFailed ⇒
  if strong_match (tag_is k) t then Failure
  else (* Now the tester knows that the ETag of [k]
        * is other than [t]. *)
        Success (data, tag_is, tag_is_not [k ↦ t::(tag_is_not k)])
| GET k t, NotModified ⇒
  if t ∈ tag_is_not then Failure
  else if (tag_is k =? Unknown) || weak_match (tag_is k) t
  then (* Now the tester knows that the ETag of [k]
        * is equal to [t]. *)
        Success (data, tag_is [k ↦ Known t], tag_is_not)
  else Failure
| GET k t0, OK t v ⇒
  if weak_match (tag_is k) t0 then Failure
  else if data k =? v
  then (* Now the tester knows the ETag of [k]. *)
        Success (data, tag_is [k ↦ Known t], tag_is_not)
  else Failure
| _, _ ⇒ Failure
end.

```

FIGURE 1.1. Ad hoc tester for HTTP/1.1 conditional requests.^a

^aPUT k t v represents a PUT request that changes k's value into v only if its ETag matches t; GET k t is a GET request for k's value only if its ETag does not match t; OK t v indicates that the request target's value is v and its ETag is t.

(i) The specification is written as a reference implementation—a nondeterministic program that exhibits all possible behavior allowed by the protocol. Inter-implementation and inter-execution uncertainties are represented by symbolic variables, and the space of nondeterministic behavior is defined by all possible assignments of the variables.

The validation logic is derived from the reference implementation, by *dualising* the server-side program into a client-side observer.

(ii) Test generation heuristics are defined as computations from the observed trace (list of sent and received messages) to the next message to send. I introduce a symbolic intermediate representation for specifying the relation between the next message and previous messages.

(iii) The symbolic language for generating test cases also enables effective shrinking of test cases. The test harness minimizes the counterexample by shrinking its symbolic representation. When running the test with a shrunk input, the symbolic representations can be re-instantiated into request messages that reflect the original heuristics.

1.3. Contribution

Symbolic abstract representation can address challenges in testing networked systems with uncertain behavior. Specifying protocols with symbolic reference implementation enables validating the system’s behavior systematically. Representing test input as abstract messages allows generating and shrinking interesting test cases. Combining these methods result in a rigorous tester that can capture protocol violations effectively.

This thesis is structured as follows:

Introduce motivation in this section entirely

For example, HTTP/1.1 servers may generate entity tags (ETags) to represent its resources’ versions [7]. The server may implement ETags with arbitrary algorithms *e.g.* hashing, timestamps *etc.*, making its value unpredictable from the clients’ requests. As a result, the tester should not compare the server’s behavior against some expectation, but evaluate whether its behavior is producible by some valid implementation.

The Deep Specification project [2] aims at building a web server and guarantee its functional correctness with respect to formal specification of the network protocol.

It is possible, but tricky, to write an ad hoc tester for HTTP/1.1 by manually “dualizing” the behaviors described by the informal specification documents (RFCs). The protocol document describes *how* a valid server should handle requests, while the tester needs to determine *what* responses received from the server are valid. For example, “If the server has revealed some resource’s ETag as **"foo"**, then it must not reject requests targetting this resource conditioned over **If-Match: "foo"**, until the resource has been modified”; and “Had the server previously rejected an **If-Match** request, it must reject the same request until its target has been modified.” Figure 1.1 shows a hand-written tester for checking this bit of ETag functionality; we hope the reader will agree that this testing logic is not straightforward to derive from the informal “server’s eye” specifications.

Networked systems are naturally concurrent, as a server can be connected with multiple clients. The network might delay packets indefinitely, so messages sent via different channels may be reordered during transmission. When the tester observes messages sent and received on the client side, it should allow all observations that can

be explained by the combination of a valid server + a reasonable network environment between the server and clients.

CHAPTER 2

Validator Theory

During the testing practice in chapter 3, the tester’s quality was evaluated by mutation testing, *i.e.* running the tester against buggy implementations to see if it rejects. To formally prove that the tester is good, I develop a theory for reasoning on testers’ good properties.

Interactive testing is a process that reveals the SUT’s interactions and determines whether it satisfies the specification. There are two kinds of interactions: (1) *inputs* that the tester can specify, and (2) *outputs* that are observed from the SUT. In particular, when testing networked systems, the input is a message sent by the tester, and the output is a message received from the SUT.

When viewing the SUT as a function from inputs to outputs, we can test the system by (1) providing an input, (2) get the output, and (3) validating the input-output pair. This process is called *synchronous testing*.

However, the nature of networked systems is that multiple messages might arrive at the system simultaneously, and a high-throughput system should handle the messages concurrently. To check the system’s validity upon concurrent inputs, the tester should send multiple messages, rather than executing “one client at a time”. This non-blocking process is called *asynchronous testing*.

My goal is to formalise the techniques in Li, Pierce, and Zdancewic [15] into a generic theory for asynchronous testing.

A tester consists of two parts: (i) a test harness that interacts with the SUT and observes the interactions, and (ii) a validator that determines whether the observations satisfy the specification.

The test harness needs to produce counterexamples effectively, and provide good coverage of test cases. The goal is to locate unknown bugs within a fixed budget, which is more practical than theoretical, and will be discussed in ???. The test theory in this dissertation focuses on guaranteeing the soundness and completeness of the validator logic.

Testers are programs that determine whether implementations are compliant or not by observing their behavior. This section defines the basic concepts and notations in testing.

DEFINITION 1 (Implementations and Behaviors). *Implementations* are programs that can interact with their environment. *Behaviors* are traces of the implementation’s interactions with the environment and consist of (i) *Outputs*, which are performed by the implementation and can be observed in the environment, and (ii) *Inputs*, which can be manipulated for testing purposes, causing (potentially) different outputs of the implementation.

“Implementation i can *produce* behavior b ” is written as “ $i \xrightarrow{b}$ ”.

DEFINITION 2 (Specifications, Validity, and Compliance). A *specification* is a description of valid behavior. “Behavior b is *valid* per specification s ” is written as “ $\text{valid}_s b$ ”.

An implementation i *complies* with a specification s (written “ $\text{comply}_s i$ ”) if it only produces behaviors that are valid per the specification:

$$\text{comply}_s i \triangleq \forall b, (i \xrightarrow{b}) \implies \text{valid}_s b$$

DEFINITION 3 (Tester components and correctness). A tester consists of (i) a *validator* that accepts or rejects implementations’ behavior, and (ii) a *test harness* that triggers different behavior with various input.

A tester is *correct* if its acceptances and rejections are sound and complete. A tester is *rejection-sound* if it only rejects in compliant implementations; it is *rejection-complete* if it can reject all in compliant implementations, provided sufficient time of execution.¹

The tester’s correctness is based on its components properties: A rejection-sound tester requires its validator to be *rejection-sound*; A rejection-complete tester consists of (i) a *rejection-complete* validator and (ii) an *exhaustive* test harness that can eventually trigger invalid behavior.

DEFINITION 4 (Soundness and completeness of validators). A validator v is *rejection-sound* with respect to specification s (written as “ $v \text{ sound}_s^{\text{rej}}$ ”) if it only rejects behaviors that are invalid per s :

$$v \text{ sound}_s^{\text{rej}} \triangleq \forall b, \neg(\text{accept}_v b) \implies \neg(\text{valid}_s b)$$

A validator v is *rejection-complete* with respect to specification s (written as “ $v \text{ complete}_s^{\text{rej}}$ ”) if it rejects all behaviors that are invalid per s :

$$v \text{ complete}_s^{\text{rej}} \triangleq \forall b, \neg(\text{valid}_s b) \implies \neg(\text{accept}_v b)$$

In property-based testing (PBT) [10], validators’ soundness and completeness are trivial, because the specification itself defines “how to compute whether some behavior is valid or not”. The validator is identical to the specification.

Whereas, in model-based testing (MBT) [4], the specification defines “how to produce valid behavior”. The validator needs to compute whether the observed behavior *can be produced* by the specification.

PBT and MBT are different views of the system under test: the former observes from outside, and the latter abstracts the internal computation. When the system might perform nondeterministic behavior, MBT allows specifying the system in a more reasonable way, which is explained in ??.

DEFINITION 5 (Exhaustiveness of test harness). A test harness h is *exhaustive* with respect to specification s (written as “ $\text{exhaustive}_s h$ ”) if, for any implementation

¹The semantics of “soundness” and “completeness” vary among contexts. This paper inherits terminologies from existing literature [17], but explicitly use “rejection-” prefix for clarity. “Rejection soundness” is equivalent to “acceptance completeness”, and vice versa.

that does not comply with the specification, the test harness can eventually trigger some invalid behavior to reveal such incompliance:

$$\begin{aligned} \text{exhaustive}_s h &\triangleq \forall i, \neg(\text{comply}_s i) \\ &\implies \exists b, (i \xrightarrow{b}_h) \wedge \neg(\text{valid}_s b) \end{aligned}$$

Exhaustive test harnesses can be built naïvely by enumerating all test cases. However, to capture bugs within realistic budget, the test harness should produce test cases that are (i) more likely to trigger bugs, and (ii) of minimal size to help analyzing and locating the bug. These challenges are further discussed in ??.

CHAPTER 3

Validator in Practice

3.1. Specification Languages

3.1.1. Property-based specification with QuickChick. My first formal specification of HTTP/1.1 was written as QuickChick [14] properties, which takes a trace of requests, and determines whether the traces is valid per protocol specification, like that shown in Figure 1.1. The specification implemented a constraint solving logic by hand, making it hard to scale when the protocol becomes more complex, as discussed in ??

3.1.2. Model-based specification with ITrees. To write specifications for protocols’ rich semantics, I employed “interaction tree” (ITree), a generic data structure for representing interactive programs, introduced by Xia et al. [19]. ITree enables specifying protocols as monadic programs that model valid implementations’ possible behavior. The model program can be interpreted into a tester program, to be discussed in Section 3.2.

Figure 3.1 defines the type `itree E R`. The definition is *coinductive*, so that it can represent potentially infinite sequences of interactions, as well as divergent behaviors. The parameter `E` is a type of *external interactions*—it defines the interface by which a computation interacts with its environment. `R` is the *result* of the computation: if the computation halts, it returns a value of type `R`.

```
CoInductive itree (E : Type → Type) (R : Type) :=
| Ret (r : R)
| Vis {X : Type} (e : E X) (k : X → itree E R)
| Tau (t : itree E R).
```

```
Inductive event (E : Type → Type) : Type :=
| Event : forall X, E X → X → event E.
```

```
Definition trace E := list (event E)
```

```
Inductive is_trace E R
: itree E R → trace E → Prop := ...
(* straightforward definition omitted *)
```

FIGURE 3.1. Interaction trees and their traces of events.

There are three ways to construct an ITree. The `Ret r` constructor corresponds to the trivial computation that halts and yields the value `r`. The `Tau t` constructor corresponds to a silent step of computation, which does something internal that does not produce any visible effect and then continues as `t`. Representing silent steps explicitly with `Tau` allows us, for example, to represent diverging computation without violating Coq’s guardedness condition [5]:

```
CoFixpoint spin {E R} : itree E R := Tau spin.
```

The final, and most interesting, way to build an ITree is with the `Vis X e k` constructor. Here, `e : E X` is a “visible” external effect (including any outputs provided by the computation to its environment) and `X` is the type of data that the environment provides in response to the event. The constructor also specifies a continuation, `k`, which produces the rest of the computation given the response from the environment. `Vis` creates branches in the interaction tree because `k` can behave differently for distinct values of type `X`.

Here is a small example that defines a type `IO` of output or input interactions, each of which works with natural numbers. It is then straightforward to define an ITree computation that loops forever, echoing each input received to the output:

```
Variant IO : Type → Type :=
| Input  : IO nat
| Output : nat → IO ().

CoInductive echo : itree IO () :=
  Vis Input (λ x ⇒ Vis (Output x) (λ _ ⇒ echo)).
```

3.2. From Specification to Tester

From an ITree specification, I conducted “offline” testing, which takes a trace and determines its validity [13], and “online” testing, where the specification is derived into a client program that validates the system under test interactively [15].

3.2.1. Offline testing of swap server. I started with testing a simple “swap server” [13], specified in Figure 3.2. The specification says that the server can either accept a connection with a new client (`obs_connect`) or else receive a message from a client over some established connection (`obs_msg_to_server c`), send back the current stored message (`obs_msg_from_server c last_msg`), and then start over with the last received message as the current state.

To test this swap server, I wrote a client program that interacts with the server and produces a trace of requests and responses, and a function that determines whether the trace `t` is a trace of the linear specification `s` *i.e.* whether `is_trace s t` in Figure 3.1 holds.

To network nondeterminism, the checker enumerates all possible server-side message orders that can explain the client-side observations, and checks if any of them satisfies the protocol specification.

```

CoFixpoint linear_spec' (conns : list connection_id)
  (last_msg : bytes) : itree specE unit :=
or ( (* Accept a new connection. *)
  c ← obs_connect;;
  linear_spec' (c :: conns) last_msg )
( (* Exchange a pair of messages on a connection. *)
  c ← choose conns;;
  msg ← obs_msg_to_server c;;
  obs_msg_from_server c last_msg;;
  linear_spec' conns msg ).

```

Definition linear_spec := linear_spec' [] zeros.

FIGURE 3.2. Linear specification of the swap server. In the `linear_spec'` loop, the parameter `conns` maintains the list of open connections, while `last_msg` holds the message received from the last client (which will be sent back to the next client). The server repeatedly chooses between accepting a new connection or doing a receive and then a send on some existing connection picked in the list `conns`. The linear specification is initialized with an empty set of connections and a message filled with zeros.

3.2.2. Online testing of HTTP. To test protocols with internal nondeterminism (*e.g.* HTTP) effectively, I introduced a symbolic representation for the server’s invisible choices, as shown in Figure A.5. I then defined a TCP network model in Figure A.4. Combining the server and network models produces a model program that exhibits all valid observations, considering both internal and network nondeterminism.

From the server and network models, I derived a tester client that interacts with servers over the network, and validates the observations against the protocol specification, as shown in Figure A.6.

Using this automatically derived tester program, I have found violations against HTTP/1.1 in the latest version of both Apache and Nginx. More details are explained in Li, Pierce, and Zdancewic [15].

3.2.3. Key innovation. To solve the problem of “determinining whether an observation is explainable by a nondeterministic program”, I reduced it into a constraint satisfiability: Although the tester doesn’t know the server and network’s exact choices, it can gain some knowledge of these invisible choices by observing the trace of messages. If the invisible choices are represented as symbolic variables, then an observed trace is valid if there exists some value for the variables that explains this trace, which can be determined by a constraint solver.

```

(* matches : (etag * exp etag) → exp bool *)
(* IF      : (exp bool * T * T) → T      *)
let put (k    : key,
        t    : etag,
        v    : value,
        data : key → value,
        xtag : key → exp etag) =
  IF (matches(t, xtag[k]),
    (* then *)
    xt := fresh_tag();
    let xtag' = update(xtag, k, xt) in
    let data' = update(data, k, v) in
    return (OK, xtag', data'),
    (* else *)
    return (PreconditionFailed, xtag, data))

```

FIGURE 3.3. Symbolic model handling conditional PUT request. The model maintains two states: `data` that maps keys to their values, and `xtag` that maps keys to symbolic variables that represent their corresponding ETags. Upon receiving a PUT request conditioned over “If-Match: `t`”, the server should decide whether the request ETag `matches` that stored in the server. Upon matching, the server processes the PUT request, and represents the updated value’s ETag as a fresh variable.

```

let tcp (buffer : list packet) =
  let absorb =
    pkt := recv();
    tcp (buffer ++ [pkt]) in
  let emit =
    let pkts = oldest_in_each_conn(buffer) in
    pkt := pick_one(pkts);
    send(pkt);
    tcp (remove(pkt, buffer)) in
  or (absorb, emit)

```

FIGURE 3.4. Network model for concurrent TCP connections. The model maintains a `buffer` of all packets en route. In each cycle, the model may nondeterministically branch to either absorb or emit a packet. Any absorbed packet is appended to the end of buffer. When emitting a packet, the model may choose a connection and send the oldest packet in it.




FIGURE 3.5. Deriving tester program from specification

CHAPTER 4

Test Harness Design

CHAPTER 5

Related Work

5.1. Specifying and Testing Protocols

Modelling languages for specifying protocols can be partitioned into three styles, according to Anand et al. [1]: (1) *Process-oriented* notations that describe the SUT's behavior in a procedural style, using various domain-specific languages like our interaction trees; (2) *State-oriented* notations that specify what behavior the SUT should exhibit in a given state, which includes variants of labelled transition systems (LTS); and (3) *Scenario-oriented* notations that describe the expected behavior from an outside observer's point of view (*i.e.*, “god’s-eye view”).

The area of model-based testing is well-studied, diverse, and difficult to navigate [1]. Here we focus on techniques that have been practiced in testing real-world programs, which includes notations (1) and (2). Notation (3) is infeasible for protocols with nontrivial nondeterminism, because the specification needs to define observer-side knowledge of the SUT's all possible internal states, making it complex to implement and hard to reason about, as shown in Figure 1.1.

Language of Temporal Ordering Specification (LOTOS) [Bolognesi1987] is the ISO standard for specifying OSI protocols. It defines distributed concurrent systems as *processes* that interact via *channels*, and represents internal nondeterminism as choices among processes.

Using a formal language strongly inspired by LOTOS, Tretmans and Laar [18] implemented a test generation tool for symbolic transition systems called TorXakis, which has been used for testing Dropbox [18].

TorXakis provides limited support for internal nondeterminism. Unlike our testing framework that incorporates symbolic evaluation, TorXakis enumerates all possible values of internally generated data, until finding a corresponding case that matches the tester's observation. This requires the server model to generate data within a reasonably small range, and thus cannot handle generic choices like HTTP entity tags, which can be arbitrary strings.

Bishop et al. [3] have developed rigorous specifications for transport-layer protocols TCP, UDP, and the Sockets API, and validated the specifications against mainstream implementations in FreeBSD, Linux, and WinXP. Their specification represents internal nondeterminism as symbolic states of the model, which is then evaluated using a special-purpose symbolic model checker. They focused on developing a post-hoc specification that matches existing systems, and wrote a separate tool for generating test cases.

5.2. Reasoning about Network Delays

For property-based testing against distributed applications like Dropbox, Hughes et al. [12] have introduced “conjectured events” to represent uploading and downloading events that nodes may perform at any time invisibly.

Sun, Xu, and Elbaum [16] symbolised the time elapsed to transmit packets from one end to another, and developed a symbolic-execution-based tester that found transmission-related bugs in Linux TFTP upon certain network delays. Their tester used a fixed trace of packets to interact with the server, and the generated test cases were the packets’ delay time.

CHAPTER 6

Discussions

CHAPTER 7

Conclusion

Bibliography

- [1] Saswat Anand et al. “An orchestrated survey of methodologies for automated software test case generation”. In: *Journal of Systems and Software* 86.8 (2013), pp. 1978–2001. ISSN: 0164-1212. DOI: <https://doi.org/10.1016/j.jss.2013.02.061>. URL: <http://www.sciencedirect.com/science/article/pii/S0164121213000563>.
- [2] Andrew W. Appel et al. “Position paper: the science of deep specification”. In: *Philosophical Transactions of the Royal Society A: Mathematical, Physical and Engineering Sciences* 375.2104 (Oct. 2017), p. 20160331. ISSN: 1364-503X. DOI: [10.1098/rsta.2016.0331](https://doi.org/10.1098/rsta.2016.0331). URL: <https://royalsocietypublishing.org/doi/10.1098/rsta.2016.0331>.
- [3] Steve Bishop et al. “Engineering with Logic: Rigorous Test-Oracle Specification and Validation for TCP/IP and the Sockets API”. In: *J. ACM* 66.1 (Dec. 2018). ISSN: 0004-5411. DOI: [10.1145/3243650](https://doi.org/10.1145/3243650). URL: <https://doi.org/10.1145/3243650>.
- [4] Manfred Broy et al. “Model-based testing of reactive systems”. In: *Volume 3472 of Springer LNCS*. Springer. 2005.
- [5] Adam Chlipala. “Infinite Data and Proofs”. In: *Certified Programming with Dependent Types*. MIT Press, 2017. URL: <http://adam.chlipala.net/cpdt/html/Cpdt.Coinductive.html>.
- [6] Roy T. Fielding and Gail Kaiser. “The Apache HTTP Server Project”. In: *IEEE Internet Computing* 1.4 (July 1997), pp. 88–90. ISSN: 1941-0131. DOI: [10.1109/4236.612229](https://doi.org/10.1109/4236.612229).
- [7] Roy T. Fielding and Julian Reschke. *Hypertext Transfer Protocol (HTTP/1.1): Conditional Requests*. RFC 7232. June 2014. DOI: [10.17487/RFC7232](https://doi.org/10.17487/RFC7232). URL: <https://rfc-editor.org/rfc/rfc7232.txt>.
- [8] Roy T. Fielding and Julian Reschke. *Hypertext Transfer Protocol (HTTP/1.1): Semantics and Content*. RFC 7231. June 2014. DOI: [10.17487/RFC7231](https://doi.org/10.17487/RFC7231). URL: <https://rfc-editor.org/rfc/rfc7231.txt>.
- [9] Roy T. Fielding et al. *Hypertext Transfer Protocol – HTTP/1.1*. RFC 2616. June 1999. DOI: [10.17487/RFC2616](https://doi.org/10.17487/RFC2616). URL: <https://rfc-editor.org/rfc/rfc2616.txt>.
- [10] George Fink and Matt Bishop. “Property-Based Testing: A New Approach to Testing for Assurance”. In: *SIGSOFT Softw. Eng. Notes* 22.4 (July 1997), pp. 74–80. ISSN: 0163-5948. DOI: [10.1145/263244.263267](https://doi.org/10.1145/263244.263267). URL: <https://doi.org/10.1145/263244.263267>.
- [11] Marijn Haverbeke. *DAV module does not respect if-unmodified-since*. Nov. 2012. URL: <https://trac.nginx.org/nginx/ticket/242>.

- [12] John Hughes et al. “Mysteries of DropBox: Property-Based Testing of a Distributed Synchronization Service”. In: *2016 IEEE International Conference on Software Testing, Verification and Validation, ICST 2016, Chicago, IL, USA, April 11-15, 2016*. 2016, pp. 135–145. DOI: 10.1109/ICST.2016.37. URL: <https://doi.org/10.1109/ICST.2016.37>.
- [13] Nicolas Koh et al. “From C to Interaction Trees: Specifying, Verifying, and Testing a Networked Server”. In: *Proceedings of the 8th ACM SIGPLAN International Conference on Certified Programs and Proofs*. CPP 2019. Cascais, Portugal: ACM, 2019, pp. 234–248. ISBN: 978-1-4503-6222-1. DOI: 10.1145/3293880.3294106. URL: <http://doi.acm.org/10.1145/3293880.3294106>.
- [14] Leonidas Lampropoulos and Benjamin C. Pierce. *QuickChick: Property-Based Testing in Coq*. Software Foundations series, volume 4. Electronic textbook, 2018. URL: <https://softwarefoundations.cis.upenn.edu/qc-current/index.html>.
- [15] Yishuai Li, Benjamin C. Pierce, and Steve Zdancewic. “Model-Based Testing of Networked Applications”. In: *ACM SIGSOFT International Symposium on Software Testing and Analysis*. 2021.
- [16] Wei Sun, Lisong Xu, and Sebastian Elbaum. “Improving the Cost-Effectiveness of Symbolic Testing Techniques for Transport Protocol Implementations under Packet Dynamics”. In: *Proceedings of the 26th ACM SIGSOFT International Symposium on Software Testing and Analysis*. ISSTA 2017. Santa Barbara, CA, USA: Association for Computing Machinery, 2017, pp. 79–89. ISBN: 9781450350761. DOI: 10.1145/3092703.3092706. URL: <https://doi.org/10.1145/3092703.3092706>.
- [17] Jan Tretmans. “Conformance testing with labelled transition systems: Implementation relations and test generation”. In: *Computer Networks and ISDN Systems* 29.1 (1996). Protocol Testing, pp. 49–79. ISSN: 0169-7552. DOI: [https://doi.org/10.1016/S0169-7552\(96\)00017-7](https://doi.org/10.1016/S0169-7552(96)00017-7). URL: <http://www.sciencedirect.com/science/article/pii/S0169755296000177>.
- [18] Jan Tretmans and Pi  re van de Laar. “Model-Based Testing with TorXakis: The Mysteries of Dropbox Revisited”. In: *Strahonja, V.(ed.), CECIIS: 30th Central European Conference on Information and Intelligent Systems, October 2-4, 2019, Varazdin, Croatia. Proceedings*. Zagreb: Faculty of Organization and Informatics, University of Zagreb. 2019, pp. 247–258.
- [19] Li-yao Xia et al. “Interaction Trees: Representing Recursive and Impure Programs in Coq”. In: *Proc. ACM Program. Lang.* 4.POPL (Dec. 2019), 51:1–51:32. ISSN: 2475-1421. DOI: 10.1145/3371119. URL: <http://doi.acm.org/10.1145/3371119>.

APPENDIX A

Unstructured contents

A.1. Challenges: Testing Internal and Network Nondeterminism

To illustrate the challenges in testing networked applications, we discuss two features of HTTP/1.1—conditional requests [7] and message forwarding [8]—showcasing internal nondeterminism and network nondeterminism, respectively.

Internal Nondeterminism. HTTP/1.1 requests can be conditional: if the client has a local copy of some resource and the copy on the server has not changed, then the server needn't resend the resource. To achieve this, an HTTP/1.1 server may generate a short string, called an “entity tag” (ETag), identifying the content of some resource, and send it to the client:

```
/* Client: */  
GET /target HTTP/1.1  
  
/* Server: */  
HTTP/1.1 200 OK  
ETag: "tag-foo"  
... content of /target ...
```

The next time the client requests the same resource, it can include the ETag in the GET request, informing the server not to send the content if its ETag still matches:

```
/* Client: */  
GET /target HTTP/1.1  
If-None-Match: "tag-foo"  
  
/* Server: */  
HTTP/1.1 304 Not Modified
```

If the tag does not match, the server responds with code 200 and the updated content as usual. Similarly, if a client wants to modify the server's resource atomically by compare-and-swap, it can include the ETag in the PUT request as `If-Match` precondition, which instructs the server to only update the content if its current ETag matches.

[LY: This is a good example, but how general is the problem, since one might question the popularity of ETags? On the other hand, if your testing framework targets application layer protocols rather than just HTTP, maybe there are more similar examples? For example, file/mail servers or databases might also require some synchronization mechanisms similar to compare-and-swap? And there might be other examples that's not compare-and-swap?][BCP: Agree that this is important to discuss.] [LYS: Mentioned at the end of this section.]

Thus, whether a server’s response should be judged *valid* or not depends on the ETag it generated when creating the resource. If the tester doesn’t know the server’s internal state (*e.g.*, before receiving any 200 response including the ETag), and cannot enumerate all of them (as ETags can be arbitrary strings), then it needs to maintain a space of all possible values, narrowing the space upon further interactions with the server.

It is possible, but tricky, to write an ad hoc tester for HTTP/1.1 by manually “dualizing” the behaviors described by the informal specification documents (RFCs). The protocol document describes *how* a valid server should handle requests, while the tester needs to determine *what* responses received from the server are valid. For example, “If the server has revealed some resource’s ETag as **“foo”**, then it must not reject requests targetting this resource conditioned over **If-Match: “foo”**, until the resource has been modified”; and “Had the server previously rejected an **If-Match** request, it must reject the same request until its target has been modified.” Figure 1.1 shows a hand-written tester for checking this bit of ETag functionality; we hope the reader will agree that this testing logic is not straightforward to derive from the informal “server’s eye” specifications.

Network Nondeterminism. When testing an HTTP/1.1 server over the network, although TCP preserves message ordering within each connection, it does not guarantee any order between different connections. Consider a proxy model in Figure A.1: it specifies how a server should forward messages. [BCP: I don’t understand why we are talking about proxies here: a simple “server + several clients” situation is enough to create network nondeterminism. (I would expect that proxying might create *additional* possibilities for nondeterminism, of course.) [LYS: We need to talk about proxy somewhere, and I didn’t find a good place elsewhere.]] [BCP: Moreover, the more I look at figures 2–5 the more confusing I find them. Only figure 5 mentions connections, but — for example, in figure 3, if we assume just a single connection between the observer and the proxy and a single connection from the proxy back to the observer, then the reordering shown in the figure is NOT valid. [LYS: Updated figure. No proxy uses the same connection for multiple requests. The proxy never knows if there’s a next request that can use the same connection.]] When the forwarded messages are scrambled as in Figure A.2, the tester should be *loose* enough to accept the server, because a valid server may exhibit such reordering due to network delays. The tester should also be *strict* enough to reject a server that behaves as Figure A.3, because no network delay can let the proxy forward a message before the observer sends it.

The kinds of nondeterminism exemplified here can be found in many other scenarios: (i) Servers may use some (unknown) algorithm to generate internal state for nonces, sequence numbers, caching metadata, *etc*, featuring internal nondeterminism. (ii) When the server runs multiple threads concurrently (*e.g.* to serve multiple clients), the operating system might schedule these threads nondeterministically. When testing the server over the network, such “nondeterminism outside the code of the server program but still within the machine on which the server is executing” is indistinguishable from nondeterminism caused by network delays, and thus can be covered by the concept “network nondeterminism.”




FIGURE A.1. Proxy model specifying a server that forwards a message immediately upon receiving it.




FIGURE A.2. A reordered observation, with two valid network-level explanations.

A.2. Specification Language

A specification in our framework consists of two parts: a server model specifying server-side behavior, [BCP: there was a discussion of this somewhere else: isn't our "application model" here just specifying HTTP and WebDAV? And so isn't it also generic? [LYS: Not generic over all L7 protocols.]] and a network model describing network delays. By composing these two models, we get a tester-side specification of valid observations over the network.

Formally, our specifications are written as *interaction trees*, a generic data structure for representing interactive programs in Coq. This language allows us to write rigorous mathematical specifications, and transform the specification into tester conveniently. In this paper, we present models as pseudocode for readability. Technical details about interaction trees can be found in [19].

Subsection A.2.1 shows how to handle network nondeterminism. Subsection A.2.2 then expands the model to address internal nondeterminism.

A.2.1. Server and Network Models. The *server model* specifies how the server code interacts with the network interface. For example, an extremely simplistic model of an HTTP proxy [BCP: again, it feels like proxies are coming out of nowhere [LYS: I'll try to make proxy more like a part of HTTP than an extension.]] (shown in Figure A.1) is written as:

```
let proxy() =
```

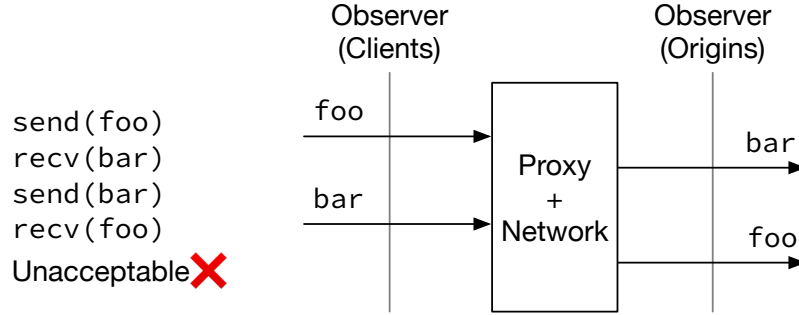


FIGURE A.3. Example of invalid observation.

```
msg := recv();
send(msg);
proxy()
```

An implementation is said to be *valid* if it is indistinguishable from the model when viewed from across the network. Consider the following proxy implementation that reorders messages: [BCP: Why are we suddenly switching to C syntax?? [LYS: To distinguish implementation from specification.]]

```
void proxy_implementation() {
    while (true) {
        recv(&msg1); recv(&msg2);
        send(msg2); send(msg1);
    }
}
```

This reordered implementation is valid, because the model itself may exhibit the same behavior when observed over the network, as shown in Figure A.2. This “implementation’s behavior is explainable by the model, considering network delays” relation is called *network refinement* by Koh et al. [13].

To specify network refinement in a testable way, we introduce the *network model*, a conceptual implementation of the transport-layer environment between the server and the tester. It models the network as a nondeterministic machine that absorbs packets and, after some time, emits them again. Figure A.4 shows the network model for concurrent TCP connections: The network either receives a packet from some node, or sends the first packet en route of some connection. This model preserves the message order within each connection, but it exhibits all possible reorderings among different connections.

The network model does not distinguish between server and tester. When one end **sends** some message, the network **recv**s the message and **sends** it after some cycles of delay; it is then observed by the other end via some **recv** call.

In Subsection A.3.3, we compose the server and network models to yield an observer-side specification for testing purposes.

A.2.2. Symbolic Representation of Nondeterministic Data. To incorporate symbolic evaluation in our testing framework, our specification needs to represent

```

let tcp (buffer : list packet) =
  let absorb =
    pkt := recv();
    tcp (buffer ++ [pkt]) in
  let emit =
    let pkts = oldest_in_each_conn(buffer) in
    pkt := pick_one(pkts);
    send(pkt);
    tcp (remove(pkt, buffer)) in
  or (absorb, emit)

```

FIGURE A.4. Network model for concurrent TCP connections. The model maintains a **buffer** of all packets en route. In each cycle, the model may nondeterministically branch to either absorb or emit a packet. Any absorbed packet is appended to the end of buffer. When emitting a packet, the model may choose a connection and send the oldest packet in it.

internally generated data as symbols. Consider HTTP PUT requests with **If-Match** preconditions: Upon success, the server generates a new ETag for the updated content, and the tester does not know the ETag’s value immediately. Our symbolic model in Figure A.5 represents the server’s generated ETags as fresh variables. The server’s future behavior might depend on whether a request’s ETag matches the generated (symbolic) ETag. Such matching produces a symbolic boolean expression, which cannot be evaluated into a boolean value without enough constraints on its variables. Our model introduces **IF** operator to condition branches over a symbolic boolean expression. Which branch the server actually took is decided by the derived tester in Section A.3.

In Subsection A.3.2, we implement the symbolic evaluation process that checks servers’ observable behavior against this symbolic model.

A.3. Derivation: from Server Specification to Testing Program

From the specified the application and network models, our framework automatically derives a tester program that interacts with the server and determines its validity. The derivation framework is shown in outline in Figure A.6. Each box is an interaction tree program, and the arrows are “interpretors” that transform one interaction tree into another. Subsection A.3.1 explains the concept of interpretation, and the rest of this section describes how to interpret the specification into a tester program.

A.3.1. Interpreting Interaction Trees. Interaction tree programs can be destructured into an interaction event followed by another interaction tree program. Such structure allows us to *interpret* one program into another. Figure A.7 shows an example of interpretation: The original **acc** program sends and receives messages, and the **tee** interpreter transforms the **acc** into another program that also prints the messages sent and received.

```

(* matches : (etag * exp etag) → exp bool *)
(* IF      : (exp bool * T * T) → T      *)
let put (k    : key,
        t    : etag,
        v    : value,
        data : key → value,
        xtag : key → exp etag) =
  IF (matches(t, xtag[k]),
      (* then *)
      xt := fresh_tag();
      let xtag' = update(xtag, k, xt) in
      let data' = update(data, k, v) in
      return (OK, xtag', data'),
      (* else *)
      return (PreconditionFailed, xtag, data))

```

FIGURE A.5. Symbolic model handling conditional PUT request. The model maintains two states: `data` that maps keys to their values, and `xtag` that maps keys to symbolic variables that represent their corresponding ETags. Upon receiving a PUT request conditioned over “If-Match: τ ”, the server should decide whether the request ETag `matches` that stored in the server. Upon matching, the server processes the PUT request, and represents the updated value’s ETag as a fresh variable.

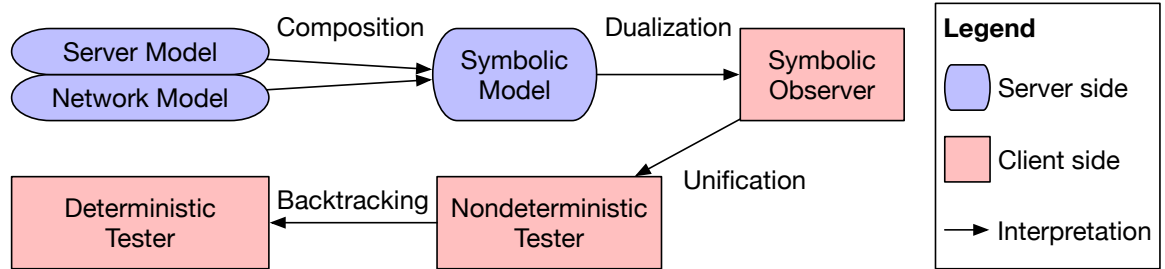


FIGURE A.6. Deriving tester program from specification

Such interpretation is done by pattern matching on the program’s structure in Line 4. Based on what the original program wants to do next, the interpreter defines what the result program should do in Line 6 and Line 8. These programs defined in accordance to events are called *handlers*. By writing different handlers for the events, interpreters can construct new programs in various ways, as shown in following subsections. Further details about interpreting interaction trees are explained by Xia et al. [19].

A.3.2. From Server Specification to Tester Program. For simplicity, we first explain how to handle servers’ internal nondeterminism with symbolic evaluation. This subsection covers a subgraph of Figure A.6, starting with dualizing the symbolic

```

let acc(sum) =                                     1
  x := recv(); send(x+sum); acc(x+sum) in          2
let tee(m) =                                       3
  match m with                                    4
  | x := recv(); m'(x) =>                          5
    a := recv(); print("IN" ++ a); tee(m'(a))      6
  | send(a); m' =>                                  7
    print("OUT" ++ a); send(a); tee(m')            8
  end in                                           9
tee(acc(0))                                       10
(* ... is equivalent to ... *)                  11
let tee_acc(sum) =                               12
  a := recv(); print("IN" ++ a);                 13
  print("OUT" ++ (a+sum)); send(a+sum);           14
  tee_acc(a+sum) in                               15
tee_acc(0)                                       16
                                              17

```

FIGURE A.7. Interpretation example. `acc` receives a number and returns the sum of numbers received so far. `tee` prints all the numbers sent and received. Interpreting `acc` with interpreter `tee` results in a program that's equivalent to `tee_acc`.

model. Here we use the server model itself as the symbolic model, assuming no reorderings by network delays. We will compose the server model with the network model in Subsection A.3.3, addressing network nondeterminism.

Dualization. To *observe* the server's behavior, we have to interpret the specified server-side events into tester-side events: When the server should send a certain message, the tester expects to receive the specified message, and rejects the server upon receiving an unexpected message; when the server should receive some message, the tester generates a message and sends it to the server, as shown in Figure A.8.

Besides sending and receiving messages, the model also has **IF** branches conditioned over symbolic expressions, like that shown in Figure A.5. Upon nondeterministic branching, the tester needs to determine which branch was actually taken, by constructing observers for both branches. Each branch represents a possible explanation of the server's behavior. Upon further interacting with the server, some branches might fail because its conjecture cannot explain what it has observed. The tester rejects the server if all branches have failed, indicating that the server corresponds to no possible case in the model.

Dualizing the server-side model produces an observer model that performs interactions to reveal the server's behavior and check its validity. This model includes all possible observations from a valid server, and needs to **determine** which branch in the server model matches the observed behavior. The model validates its observations with unification events **unify** and **guard**. These primitive events are handled by later interpretations: The **unify** and **guard** events in each branch are instantiated into

```

let observe (server) =
  match server with
  | pkt := recv(); s'(pkt) =>
    p := gen_pkt(); send(p); observe (s'(p))
  | send(pkt); s' =>
    p := recv(); guard(pkt, p); observe (s')
  | IF (x, s1, s2) =>
    (* Allow validating observation with [s1],
     * provided [x] is unifiable with [true];
     * Or, unify [x] with [false],
     * and validate observation with [s2]. *)
    determine(unify(x, true ); observe (s1),
              unify(x, false); observe (s2))
  | r  := _(); s'(r) =>
    r1 := _(); observe (s'(r1))
end

```

FIGURE A.8. Dualizing server model into observer model. Upon **recv** events, the observer generates a packet and sends it to the server. For **send** events, the observer receives a packet **p1**, and fails if it does not match the specified **pkt**. When the server makes nondeterministic **IF** branches, the observer **determine**s between the branches by **unify**ing the branch condition with its conjectured value, and then observing the corresponding branch.

symbolic evaluation logic that decides whether this branch should fail or not; The **determine** events are instantiated into backtracking searches to find if all branches have failed, which rejects the server.

Symbolic Evaluation. In this interpretation phase, we handle nondeterminism at data level by handling **fresh** events in the server model, as well as **unify** and **guard** events introduced by dualization. The interpreter instantiates these events into symbolic evaluation algorithms.

As shown in Figure A.9 (skip Line 18–28 for now—we’ll explain that part later), the tester checks whether the observed/conjectured value matches the specification, by maintaining the constraints on the symbolic variables. These constraints are initially empty when the variables are generated by **fresh** events. As the test runs into **unify** and **guard** events, it adds constraints **asserting** that the observed value matches the specification, and checks whether the constraints are still compatible. Incompatibility among constraints indicates that the server has exhibited behavior that cannot be explained by the model, implying violation against the current branch of specification.

Handling Incoming Connections. In addition to generating data internally, the server might exhibit another kind of nondeterminism related to the outgoing connections it creates. For example, when a client uses an HTTP server as proxy, requesting resources from another server, the proxy server should create a new connection to the

```

(* unifyS = list variable * list constraint *)
(* new_var : unifyS → variable * unifyS *)
(* assert : exp T * T * unifyS → option unifyS *)
let unifier (observer, map : mcid → pcid,
            vars : unifyS) =
  match observer with
  | x := fresh(); o'(x) ⇒
    let (x1, vars') = new_var(vars) in
    unifier (o'(x1), vars', map)
  | unify(x, v); o' ⇒
    match assert(x, v, vars) with
    | Some vars' ⇒ unifier (o', vars', map)
    | None ⇒ failwith "Unexpected payload"
    end
  | guard(p0, p1); o' ⇒
    match assert(p0, p1, vars) with
    | Some vars' ⇒
      let mc = p0.source in
      let pc = p1.source in
      if mc.is_created_by_server
      then match map[mc] with
        | pc ⇒ unifier (o', vars', map)
        | unknown ⇒
          let map' = update(map, mc, pc) in
          unifier (o', vars', map')
        | others ⇒
          failwith "Unexpected connection"
          end
      else unifier (o', vars', map)
    | None ⇒ failwith "Unexpected payload"
    end
  | r := _(); o'(r) ⇒
    r1 := _(); unifier (o'(r1), vars, map)
  end
end

```

FIGURE A.9. Instantiating symbolic events. The tester maintains a `unifyState` which stores the constraints on symbolic variables. When the specification creates a `fresh` symbol, the tester creates an entry for the symbol with no initial constraints. Upon `unify` and `guard` events, the tester checks whether the `assert` is compatible with the current constraints. If yes, it updates the constraints and move on; otherwise, it raises an error on the current branch.

target server. However, as shown in Figure A.2, when the tester receives a request from an accepted connection, it does not know which client’s request the proxy was forwarding, due to network delays.

Outgoing connections created by the server model are identified by “model connection identifiers” (`mcid`), and the tester accepts incoming connections identified by “physical connection identifiers” (`pcid`). As shown in Line 18–28 of Figure A.9, to determine which `mcid` in the specification does a runtime `pcid` corresponds to, the tester maintains a mapping between the connection identifiers. Such mapping ensures the tester to check interactions on an accepted connection against the right connection specified by the server model.

Backtracking. Symbolic evaluation determines whether the observations matches the tester’s conjectures on each branch. So far, the derived tester is a nondeterministic program that rejects the server if and only if all possible branches have raised some error. To simulate this tester on a deterministic machine, we execute one branch until it fails. Upon failure in the current branch, the simulator switches to another possible branch, until it exhausts all possibilities and rejects the server, as shown in Line 9–13 of Figure A.10.

When switching from one branch to another, the tester cannot revert its previous interactions with the server. Therefore, it must match the server model against all interactions it has performed, and filter out the mismatching branches, as shown in Line 15 and Line 21 of Figure A.10.

We’ve now derived a tester from the server model. The specified server runs forever, and so does the tester (upon no violations observed). We accept the server if the tester hasn’t rejected it after some large, pre-determined number of steps of execution.

Test Case Generation. Counterexamples are sparsely distributed, especially when the bugs are related to server’s internally generated data like ETags, which can hardly be matched by a random test case generator. After observing the `ETag` field of some response, the generator can send more requests with the same `ETag` value, rather than choosing an unknown value arbitrarily.

As shown in Figure A.11, our derivation framework allows passing the programs’ internal state as the events’ parameters, so the test case generator can utilize the states in all intermediate interpretation phases, and apply heuristics to emphasise certain bug patterns.

Notice that the state-passing strategy only allows tuning *what* messages to send. To reveal bugs more efficiently in an interactive scenario, we need to tune *when* the interactions are made, which is further discussed in Subsection A.4.2. Generating test cases in certain orders is to be explored in future work.


```

(* filter : event T * T * list M → list M *)
(* [filter(e, r, l)] returns a subset in [l],
* where the model programs' next event is [e]
* that returns [r]. *)
let backtrack (current, others) =
  match current with
  | determine(t1, t2) ⇒
    backtrack (t1, t2::others)
  | failwith error ⇒ (* current branch failed *)
    match others with
    | [] ⇒ failwith error
    | another::ot' ⇒ backtrack (another, ot')
  end
  | send(pkt); t' ⇒
    let ot' = filter(SEND, pkt, others) in
    send(pkt); backtrack (t', ot')
  | pkt := recv(); t'(pkt) ⇒
    opkt := maybe_recv();
    match opkt with
    | Some p1 ⇒
      let ot' = filter(RECV, pkt, others) in
      backtrack (t'(p1), ot')
    | None ⇒
      (* no packet arrived *)
      match others with
      | [] ⇒ backtrack (current, []) (* retry *)
      | another::ot' ⇒
        (* postpone *)
        backtrack (another, ot'++[current])
      end
    end
  end
end in
backtrack (tester_nondet, [])

```

FIGURE A.10. From nondeterministic model to deterministic tester program. If the model makes nondeterministic branches, the tester picks a branch to start with, and puts the other branch into a set of other possibilities. If the current branch has failed, the tester looks for other possible branches to continue checking. When the current branch sends a packet, the tester filters the set of other possibilities, and only keeps the branches that match the current send event. If the model wants to receive a packet, the tester handles both cases whether some packet has arrived or not.

```

let http_server (http_st) =
  request := recv_HTTP(http_st);
  (response, st') := process(request, http_st);
  http_server (st')
...
let observer (server) =
  match server with
  | req := recv_HTTP(http_st); s'(req) ⇒
    r1 := gen_Observer(http_st);
    send(r1); observe (s'(r1))
...
let unifier (observer, vars, conn) =
  match observer with
  | req := gen_Observer(http_st); o'(req) ⇒
    r1 := gen_Unifier(http_st, vars, conn);
    unifier (o'(r1), vars, conn)
...

```

FIGURE A.11. Embedding programs' internal state into the events. By expanding the events' parameters, we enrich the test case generator's knowledge along the interpretations.

A.3.3. Network Composition. We have shown how to derive a tester from the server model itself. The server model describes how a reference server processes messages. For protocols like HTTP/1.1 where servers are expected to handle one request at a time, a reasonable server model should be “linear” that serves one client after another. As a result, the derived tester only simulates a single client, and does not attempt to observe the server's behavior via multiple simultaneous connections.

The network model describes how messages sent by one end of the network are eventually received by the other end. When interacting with multiple clients, a valid server's observable behavior should be explainable by “server delayed by the network”, as discussed in Subsection A.2.1. To model this set of observations, we compose the server and network models by attaching the server model as one end on the network model.

As shown in Figure A.12, we **compose** the events of server and network models. Messages sent by the server are received by the network and sent to clients after some delay, and vice versa. Such composition produces a model that branches nondeterministically, and includes all possible interactions of a valid HTTP server that appear on the client side.

The composed model does not introduce new events that were not included in the server model: The network model in Figure A.4 does perform nondeterministic **or** branches, but **or**(*x*,*y*) is a syntactic sugar for **b** := **fresh**(); **IF**(*b*,*x*,*y*). Therefore, using the same derivation algorithm from the server model to single-connection tester

```

1  let compose (net, bi, bo, srv) =
2    let step_net =
3      match net with
4      | send(pkt); n' =>
5        if pkt.to_server
6        then compose (n', bi++[pkt], bo, srv)
7        else send(pkt);  (* to client *)
8          compose (n', bi, bo, srv)
9      end
10     | pkt := recv(); n'(pkt) =>
11       match bo with
12       | p0::b' => compose (n'(p0), bi, b', srv)
13       | []      => p1 := recv();
14                   compose (n'(p1), bi, bo, srv)
15       end
16     | r := _(); n'(r) =>
17       r1 := _(); compose (n'(r1), bi, bo, srv)
18     end in
19   match srv with
20   | send(pkt); s' =>
21     compose (net, bi, bo++[pkt], s')
22   | pkt := recv(); s'(pkt) =>
23     match bi with
24     | p0::b' => compose (net, b', bo, s'(p0))
25     | []      => step_net
26     end
27   | r := _(); s'(r) =>
28     r1 := _(); compose (net, bi, bo, s'(r1))
29   end in
30   compose (tcp, [], [], http)
31

```

FIGURE A.12. Composing `http` server model with `tcp` network model by interpreting their events and passing messages from one model to another. The composing function takes four parameters: server and network models as `srv` and `net`, and the message buffers between them. When `srv` wants to `send` a packet in Line 21, the packet is appended to the outgoing buffer `bo` until absorbed by `net` in Line 12, and eventually emitted to the client in Line 7. Conversely, packets sent by clients are absorbed by `net` in Line 13, emitted to the application's incoming buffer `bi` in Line 6, until `srv` consumes it in Line 24.

program, we can derive the composed server+network model into a multi-connection tester.

Notice that the server and network events are scheduled at different priorities: The composition algorithm steps into the network model lazily, not until the server is blocked in Line 25. When the network wants to `recv` some packet in Line 10, it prioritizes packets sent by the server, and only receives from the clients if the server’s outgoing buffer has been exhausted. Such design is to enforce the tester to terminate upon observing invalid behavior: When the server’s behavior violates the model, the tester should check all possible branches and determine that none of them can lead to such behavior. If the model steps further into the network, it would include infinitely many `absorb` branches in Figure A.4, so the derived tester will never exhaust “all” branches and reject the server. Scheduling network events only when the server model is blocked produces sufficient nondeterminism to accept valid servers.

A.4. Evaluation

To evaluate whether our derived tester is effective at finding bugs, we ran the tester against mainstream HTTP servers, as well as server implementations with bugs inserted by us.

A.4.1. Experiment Setup.

Systems Under Test (SUTs). We ran the tests against Apache HTTP Server [6], which is among the most popular servers on the World Wide Web. We used the latest release 2.4.46, and edited the configuration file to enable WebDAV and proxy modules. Our tester found a violation against RFC 7232 in the Apache server, so we modified its source code before creating mutants.

We’ve also tried testing Nginx and found another violation against RFC 7232. However, the module structure of Nginx made it difficult to fix the bug instantly. (The issue was first reported 8 years ago and still not fixed!) Therefore, no mutation testing was performed on Nginx.

Infrastructure. The tests were performed on a laptop computer (with Intel Core i7 CPU at 3.1 GHz, 16GB LPDDR3 memory at 2133MHz, and macOS 10.15.7). The SUT was deployed as a Docker instance, using the same host machine as the tester runs on. They communicate with POSIX system calls, in the same way as over Internet except using address `localhost`. The round-trip time (RTT) of local loopback is 0.08 ± 0.04 microsecond (at 90% confidence).

A.4.2. Results.

Finding Bugs in Real-World Servers and Mutants. Our tester rejected the unmodified Apache HTTP Server, which uses strong comparison for PUT requests conditioned over `If-None-Match`, while RFC 7232 specified that `If-None-Match` preconditions must be evaluated with weak comparison[BCP: What are strong and weak comparison? [LYS: ETag jargons.]]. We reported this bug to the developers, and figured out that Apache was conforming with an obsoleted HTTP/1.1 standard [9]. The latest standard has changed the semantics of `If-None-Match` preconditions, but Apache didn’t update the logic correspondingly.




FIGURE A.13. Cost of detecting bug in each server/mutant. The left box with median line is the tester’s execution time before rejecting the server, which includes interacting with the server and checking its responses. The right bar with median circle is the number of HTTP/1.1 messages sent and received by the tester before finding the bug. Results beyond 25%–75% are covered by whiskers.

We created 20 mutants by manually modifying the Apache source code. The tester rejected all the 20 mutants, located in various modules of the Apache server: `core`, `http`, `dav`, and `proxy`. They appear both in control flow (*e.g.*, early return, skipped condition) and in data values (*e.g.*, wrong arguments, flip bit, buffer off by one byte).

We didn’t use automatic mutant generators because (i) Existing tools could not mutate all modules we’re interested in; and (ii) The automatically generated mutants could not cause semantic violations against our protocol specification.

When testing Nginx, we found that the server did not check the preconditions of PUT requests. We then browsed the Nginx bug tracker and found a similar ticket opened by Haverbeke [11]. These results show that our tester is capable of finding bugs in server implementations, including those we’re unaware of.

Performance. As shown in Figure A.13, the tester rejected all buggy implementations within 1 minute. In most cases, the tester could find the bug within 1 second.

Some bugs took longer time to find, and they usually required more interactions to reveal. This may be caused by (1) The counter-example has a certain pattern that our generator didn’t optimize for, or (2) The tester did produce a counter-example, but failed to reject the wrong behavior. We determine the real cause by analysing the bugs and their counterexamples:

- Mutants 19 and 20 are related to the WebDAV module, which handles PUT requests that modify the target’s contents. The buggy servers wrote to a different target from that requested, but responds a successful status to the

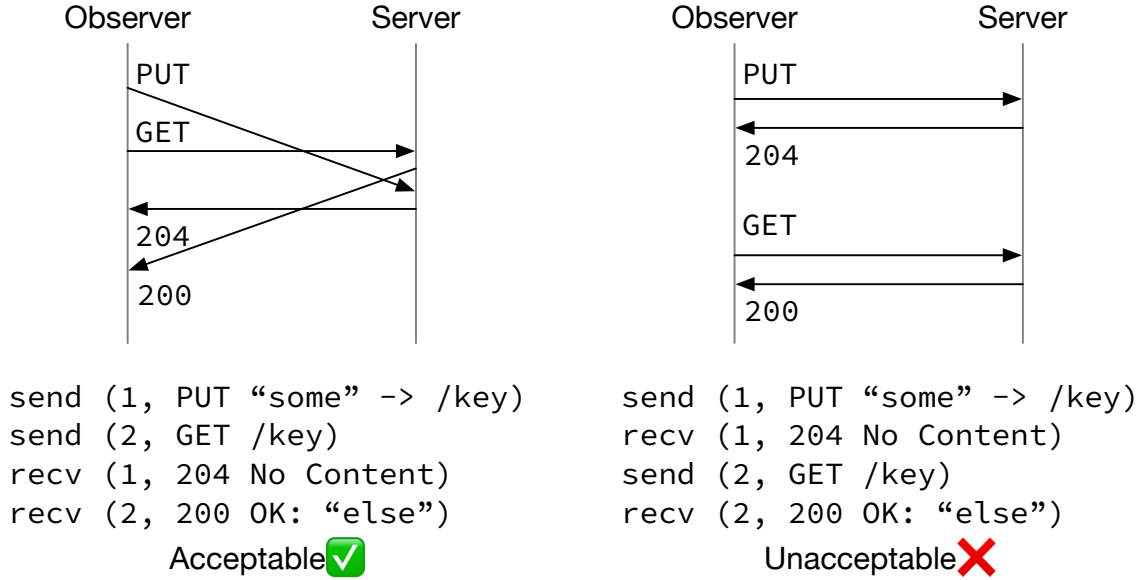


FIGURE A.14. The trace on the left does not convince the tester that the server is buggy, because there exists a certain network delay that explains why the PUT request was not reflected in the 200 response. When the trace is ordered as shown on the right, the tester cannot imagine any network reordering that causes such observation, thus must reject the server.

client. The tester cannot tell that the server is faulty until it queries the target's latest contents and observes an unexpected value. To reject the server with full confidence, these observations must be made in a certain order, as shown in Figure A.14.

- Mutant 18 is similar to the bug in vanilla Apache: the server should have responded with 304 Not Modified, but sent back 200 OK instead. To reveal such violation, a minimal counterexample consists of 4 messages: (1) GET request, (2) 200 OK response with some ETag x , (3) GET request conditioned over `If-None-Match: x` , and (4) 200 OK response, indicating that the ETag x did not match itself. Notice that (2) must be observed before (3), otherwise the tester will not reject the server, with a similar reason as Figure A.14.
- Mutant 5 causes the server to skip some code in the core module, and send nonsense messages when it should respond with 404 Not Found. The counterexample can be as small as one GET request on a non-existential target, followed by a non-404, non-200 response. However, our tester generates request targets within a small range, so the requests' targets are likely to be created by the tester's previous PUT requests. Narrowing the range of test case generation might improve the performance in aforementioned Mutants 18–20, but Mutant 5 shows that it could also degrade the performance of finding some bugs.

- The mutants in proxy module caused the server to forward wrong requests or responses. When the origin server part of the tester accepts a connection from the proxy, it does not know for which client the proxy is forwarding requests. Therefore, the tester needs to check the requests sent by all clients, and make sure none of them matches the incoming proxy request, before rejecting the proxy.

These examples show that the time-consuming issue of some mutants are likely caused by limitations in the test case generators. Cases like Mutant 5 can be optimized by tuning the request generator based on the tester model's runtime state, but for Mutants 18–20, the requests should be sent at specific time periods so that the resulting trace is unacceptable per specification. How to produce a specific order of messages is to be explored in future work.